

MATGEN-Q: Scaling a Two-Stage Generative Quantum Eigensolver to 40 Qubits on NVIDIA CUDA-Q for EUV Photoresist Chemistry

Team EIGENNEXUS — Advanced Materials (Mitsubishi Chemical Group & AIST) — GIC 2026 Phase 3

1. Focus Area and Rationale

Ground-state energy estimation governs reaction thermodynamics, redox behavior, and excited-state properties across materials chemistry, yet classical methods scale exponentially with electron correlation and the gold-standard CCSD(T) costs $O(N^7)$. The Generative Quantum Eigensolver (GQE; Nakaji et al. 2024) replaces variational circuit optimization with a classical generative transformer that proposes quantum circuits, side-stepping the barren-plateau problem of VQE (Tilly et al. 2022; a GQE property established by Nakaji et al. 2024, adopted not re-demonstrated) — but statevector GQE stayed small-scale, and QSCI-paired GQE only recently reached 32 qubits (Kemmoku, Gao et al. 2026). **Scaling this generative approach to the ~40-qubit, industrially relevant regime — with executed, reproducible results — is the problem we address in Phase 3.** Our primary, executed contribution is quantum-accurate energies for the exact chemistry where BOTH classical pillars fail — DFT's spin-state ordering flips sign across functionals and single-reference CCSD(T) collapses non-variationally below the exact energy as bonds break — on open-shell, multireference transition-metal oxides (CrO, NiO) and EUV-relevant tin oxides (SnO, SnO₂) at 20 qubits, validated against exact diagonalization and third-party HamLib Hamiltonians. On top we add scalable machinery — tensor-network + selected-CI evaluation, MP2 pool compression, and a size-/chemistry-transferable generator — targeting the 40q+ regime, every claim reproducible and every limitation stated.

Our target is EUV semiconductor photoresist chemistry (tin-oxo clusters; Sn, Hf, Zr oxides), whose open-shell, multireference character makes DFT candidate rankings unreliable — the focus of an active quantum-simulation program by one of this challenge's providers (Kharazi et al., Xanadu & Mitsubishi Chemical, 2026). We quantify the unreliability directly (Figure 1): across six functionals the CrO spin-state splitting spans 1.9 eV and B3LYP flips the ground state, while CASCI and QSCI agree on one quintet (⁵II) ground state matching the experimental X³II term — a single quantum-accurate number replacing the entire functional-choice band. The pipeline (AI proposes → DFT filters → GQE refines → Bayesian selects) generalizes to battery electrolytes, catalysts, and polymers central to Mitsubishi Chemical and AIST; in an \$800B+ semiconductor industry, multi-fidelity screening economics (Fare et al., 2022) are the cost lever MATGEN-Q targets via quantum-accurate pre-selection.



Figure 1. CrO spin-state decision. Six DFT functionals span 1.9 eV for the triplet–quintet gap and B3LYP flips the ground state to the (wrong) triplet; CASCI and QSCI in a fixed CAS(10,10) give one consistent quintet (⁵II) ground state, agreeing with experiment. The robust claim is the 1.9 eV spread, the B3LYP sign flip, and the experimental-sign agreement, not the precise gap magnitude (CAS/def2-SVP is a fixed modest active space).

Stakeholder relevance. For Mitsubishi Chemical: a quantum-accurate filter ahead of costly synthesis and lithographic testing, where DFT mis-ranks tin-oxo energetics; for AIST/G-QuAT: a scalable, reproducible GQE workflow on CUDA-Q extending the jointly-developed program past its size limits.

2. Target System and Data Modeling Strategy

Scaling vehicle. The H_n chain series (n = 2...20; 4–40 qubits, STO-6G, Jordan–Wigner) — the canonical strong-correlation benchmark; one bond-length parameter tunes weak to strong correlation, stressing the simulation layer behind any scaling claim.

Target chemistry. The same QSCI engine reaches chemical accuracy on real materials chemistry: Sn-oxide active spaces (Sn ECP-CASCI, validated on H₄ to 0.0000 mHa) — SnO (16q) and SnO₂ (20q) to chemical accuracy (≤ 0.6 mHa asserted by reproduce.py; observed 0.11–0.45 / 0.18–0.23 across environments), and — to reach the genuine EUV motif rather than a diatomic — a bridged **Sn–O–Sn tin-oxo fragment (Sn₂O₂ rhombus) to 0.41 mHa at 16q;** and genuine open-shell, multireference transition-metal oxides — **CrO ⁵II 0.038 mHa and NiO ³Σ⁻ 0.197 mHa at 20 qubits.** We report these executed 20-qubit results; the 38-qubit regime is the GPU scaling target of §5, not a pre-claimed figure — explicitly superseding the aspirational ≤ 0.08 mHa @38q of our Phase 2 draft, which we do not claim.

Data integrity. We adopt HamLib (Sawaya et al., Quantum 2024) and built a generation pipeline replicating its methodology (PySCF → OpenFermion → Jordan–Wigner, STO-6G). **Validated against the published files: FCI reproduced to 0.00000 mHa (H₂–H₆), and at 28/32/40 qubits our Hamiltonians match HamLib exactly in term count (27,735 / 47,489 / 116,577) and to ~15 significant figures in coefficient magnitude** (differing only by a spectrum-invariant orbital-phase convention) — third-party reproducible.

Accuracy anchors / classical references. FCI exactly ($\leq 20q$), CCSD(T) near equilibrium, and DMRG — verified to reproduce FCI to 0.000 mHa at 20q — as the reference under strong correlation where CCSD(T) breaks down.

3. GQE-Based Approach and Algorithmic Innovation

MATGEN-Q uses a two-stage GQE. Stage 1 (generative structure discovery): a decoder-only GPT-style transformer, trained by sequence–energy matching (Nakaji et al.), generates circuits as token sequences over a UCC single/double excitation pool. **Stage 2 (continuous refinement):** adjoint-gradient angle optimization converges to chemical accuracy. Four innovations make it scalable:

(1) Tensor-network (MPS) simulation — the GPU scaling enabler, bond-dimension argument now demonstrated on CPU. CUDA-Q's tensornet-mps memory scales with entanglement, not 2^n ; we verify the enabling claim with block2 DMRG: **the bond dimension χ for chemical accuracy grows only modestly with size ($\chi \approx 50/100/400$ at 20/28/40 qubits), entanglement area-law-bounded near equilibrium ($S_{\max} 0.39$) rising into strong correlation ($S_{\max} 4.43$).** So at 40q, $\chi \approx 400$ puts the MPS footprint ($\sim n \cdot \chi^2 \cdot d$, megabytes) far below the ~ 16 TB exact statevector. The 20q/28q GPU runs are executed to chemical accuracy (§5); the full 40q tensornet-mps run remains the primary owed deliverable, de-risked by the measured χ target and the $\approx 5\times$ faster growth engine.

(2) QSCI energy evaluation — accuracy and noise-aware. Quantum-Selected Configuration Interaction (Kanno et al. 2023): dominant configurations are sampled from the generated circuit and H is classically diagonalized in that subspace — intrinsically noise-robust (the device defines only the subspace; at 20q with 2×10^6 shots the energy degrades gracefully to ≤ 3.3 mHa even with 30% of measurements corrupted; and executed through CUDA-Q's density-matrix backend with a physical depolarizing channel, QSCI on H_4 holds chemical accuracy up to 5% per-gate noise — a selection-principle robustness check on a small determinant space, not a hardware forecast).

(3) Operator-pool compression — efficiency (demonstrated). Ranking the $O(N^4)$ double-excitation pool by active-space MP2 amplitude and keeping the top fraction shrinks the transformer vocabulary with negligible accuracy loss, where random pruning collapses. Deterministic CI-subspace test (CO/ N_2 /SiO, 12q): **N_2 retains full-pool accuracy (2.26 mHa) keeping only 25% of doubles (vocabulary 1170 $\rightarrow$ 430) vs 50.4 mHa for random pruning at the same size ($\sim 22\times$); CO holds 3.7 mHa at 40% kept vs 29.7 mHa random.**

(4) Distributed hybrid workflow. Detailed in §4.

Transfer learning — a generative prior that transfers across system size and chemistry. Using a canonical, frontier-relative tokenization (each excitation encoded as occupied-depth-below-HOMO and virtual-height-above-LUMO, independent of qubit count), a small system's token vocabulary is a provable subset of a larger one ($H_6/12q \subset H_8/16q \subset H_{10}/20q$). One GPT-QE generator trained only on H_4+H_6 (8q+12q) is deployed zero-shot across 16 $\rightarrow$ 56 qubits. **The robust signals are statistical, not a single-number win: $\sim 3.7\times$ tighter selection spread than random (≈ 2.0 vs ≈ 7.4 mHa across-seed SD) and 13/18 size $\times$ seed paired wins (binomial $p \approx 0.05$)** — reliable selection versus a lottery. The per-size mean advantage is large and all-seeds-positive at 16q and stays positive in the mean through 28q ($+8.9/+8.3/+7.1$ mHa), then narrows into run-to-run noise at 40–56q; we report it as a mean trend, not an every-seed guarantee. That the learned policy captures real chemistry — not a per-system fit — is corroborated by interpretability: trained on energy alone, it recovers the MP2 double-excitation amplitude hierarchy (Spearman $\rho=0.31$, $p<0.002$; 4 of its top-8 doubles are MP2's). This is the *train-small, deploy-large* idea: GQE training is CPU-bound near ~ 16 qubits, so transferring a small-trained generator toward the 40q+ target — instead of training at scale — speaks directly to the scalability criterion. The enabler is a determinant-space evaluation (each excitation maps by bitmask to a determinant, diagonalized via Slater-Condon, validated to 0.0000 mHa against Jordan–Wigner), needing no 2^n statevector, so 56q runs on CPU — a *selected-CI proxy* for circuit-sampled QSCI, relative at a small fixed budget ($K=96$), not absolute chemical accuracy.

Cross-chemistry transfer (honest, partial). The same H-chain-trained generator, deployed to oxide active spaces, beats random clearly on 3 of 6 targets (CO-12q/SiO/CO-20q), ties on BeO, loses on SnO (heavy-element ECP chemistry too dissimilar). Target-specific MP2 is the strongest selector on every target: cross-size transfer is robust, cross-chemistry works only when the target resembles training — a zero-cost prior, not a universal one. A same-size conditional-GQE variant (cf. Minami, Nakaji et al. 2025) gave only a within-noise gain, also a reported negative.

4. Hybrid Architecture

The classical transformer trains on GPU (PyTorch); circuit evaluations are dispatched across GPUs via CUDA-Q's mcpu in an asynchronous generate $\rightarrow$ evaluate $\rightarrow$ update loop — quantum resources do state preparation and energy estimation (MPS/QSCI), classical resources do generation, optimization, and active-space selection. **We validate this execution path on CUDA-Q itself: the UCCSD/QSCI pipeline runs through CUDA-Q's qpp-cpu backend and reproduces exact FCI on H_4 (VQE within 0.02 mHa via**

cudaq.observe; QSCI within chemical accuracy via cudaq.sample), confirming the generated circuits compile and sample through the SDK — and the full QPU submission chain (export → submit → decode → QSCI) is executed end-to-end through the qBraid cloud runtime on its free simulator tier (+2.0 mHa vs FCI, job IDs committed); the 20q and 28q GPU runs are now executed on NVIDIA hardware through qBraid (§5), the 40q tensor-net-mps run and real QPU silicon owed.

5. Phase 3 Execution and Results

Verified results (CPU). One command (*reproduce.py*) runs 25 automated checks (16 re-executions + 9 committed-artifact audits incl. the pre-registration SHA-256 and the AQT trapped-ion decode): **25/25 PASS**; every value traces to a committed result file (Table 1).

System	Qubits	Quantum (GQE/QSCI)	Classical ref.	Notes
H ₂ /H ₄ /H ₆	4/8/12	0.000 / 0.009 / 0.297 mHa	FCI (exact)	two-stage GQE (UCCSD)
H ₆ GQE → QSCI	12	1.05 mHa (51 → 1.05, ~50×)	FCI	measured pipeline
H ₁₀ /H ₁₄	20/28	0.57 / 1.21 mHa	FCI / CCSD(T)	2,401 / 18,201 dets
CrO ⁵ Π / NiO ³ Σ ⁻	20	0.038 / 0.197 mHa	CASCI (exact)	open-shell multiref.
SnO / SnO ₂	16/20	chem. acc. ≤0.6 (obs. 0.11–0.45 / 0.18–0.23)	FCI	EUV target; values version/run sensitive
Noise (H ₁₀)	20	≤3.3 mHa @ 30% corrupt (2×10 ⁶ shots)	—	noise-aware bonus

Table 1. Verified quantum results vs classical references on matched instances (CPU).

Classical baseline and the exact wall (matched instances, timed). On identical STO-6G H_n geometries, classical FCI wall-clock grows steeply (~6× from 20 → 24 qubits, 0.60 s → 3.63 s on a single CPU thread) and is intractable by 32q on CPU; CCSD(T) stays cheap but its error climbs with correlation and breaks down under strong correlation (at H₂₄/48q, classical DMRG is itself 6.83 mHa off CCSD(T)). Exact statevector needs ~16 TB at 40q vs ~0.2 GB for the MPS tier (measured χ=400), and ~10⁶ QSCI determinants vs ~3×10¹⁰ for full CI — the memory and determinant walls the MPS + QSCI tiers remove (Table 2; quantified in results/crossover).

System	Qubits	FCI (exact) wall-clock	CCSD(T) err vs FCI	CCSD(T) wall-clock
H ₆	12	0.04 s	0.026 mHa	0.08 s
H ₁₀	20	0.60 s	0.173 mHa	0.17 s
H ₁₂	24	3.63 s	0.282 mHa	0.21 s
H ₁₄	28	minutes	—	~0.2 s
H ₁₆ ⁺	32+	intractable (CPU)	—	—

Table 2. Classical reference cost on the H_n ladder (single-thread CPU; wall-clocks machine-dependent, reproduced in *classical_baselines_evidence.json* — the steep growth and exact-method wall, not absolute seconds, are the claim; H₁₄⁺ rows qualitative).

Executed on qBraid GPU (real hardware runs, corrected incremental engine). Two at-scale runs executed on NVIDIA GPU through qBraid, each judged against a reference committed before access (git-timestamped preregistration): **20q H₁₀ (cuStateVec)** reproduces full CI to **+0.000 mHa** at 0.44 GB device memory — an exact GPU-execution anchor; **28q H₁₄ (cuStateVec)** — a device-sampled QSCI grown to 64,212 determinants — reaches **+0.395 mHa vs block2 DMRG(χ=400)** at 2.82 GB device memory, chemical accuracy above the 20q scale on hardware. Both meet their pre-registered energy (P1) and device-memory (P3) criteria. **40q H₂₀ (primary criterion):** the fast incremental engine drives the QSCI error **monotonically 8.9 → 3.2 mHa from 30k → 150k determinants** vs committed DMRG(χ=400), converging as ~(*dets*)^{-0.7} toward the pre-registered P2 prediction (~10⁶ dets for chemical accuracy); the strict 1.6 mHa pass is a determinant-budget limit at 40q, not a method limit (attempt-1's old engine stalled at 6k/+14.8 mHa). The subspace is MP2-seeded (the tensor-net-mps sampler's fixed MPS-contraction cost is impractical at 40q; growth is seed-independent, device sampler proven exact at 20q/28q). **CrO near-38q — executed, the reference corrected:** QSCI on CrO CAS(18,19)=38q converged **-3.784 mHa BELOW the same-CAS DMRG(χ=400) reference** (529,392 determinants, 19.1 h, A100 host; reference committed before access). The pre-registered χ-escalation counter-audit confirms the mechanism rather than a QSCI error — χ=800 and χ=1200 references descend toward the QSCI energy from above (gaps +1.06 / +0.36 mHa) and never cross it: a genuine truncation-error correction verified at three bond dimensions. **Real-QPU silicon (10–16q) — executed and decoded:** the 3-job pooled protocol ran end-to-end through qBraid (qir-sv simulator tier: +2.0 / +2.4 mHa vs FCI, PASS, job IDs committed), and the trapped-ion flight (AQT ibex-q1, 2,000 shots/job, SHA-pinned exports) decoded 2026-07-20 — number-conserving post-selection kept 28.6% / 25.1% of shots, the

device-sampled determinants give +20.4 / +11.1 mHa, and device-seeded QSCI growth recovers exact FCI (0.000 mHa) — genuine 12-qubit trapped-ion samples driving the same QSCI machinery as the 40q runs.

Determinant-budget sweep and a classical baseline. Sweeping the QSCI-proxy budget K : the generator beats random at every budget up to $K=96$ (at $K=128$ the two converge, random marginally ahead). Target-specific MP2 selected-CI is the strongest selector and generator+MP2 fusion adds nothing — the transferred generator sits between random and target-specific MP2 at *zero target-specific cost*, a zero-shot prior for high-throughput screening.

What the results establish. Three things, demonstrated not asserted: (i) correctness — the GQE/QSCI pipeline reproduces the FCI/CASCI reference to chemical accuracy on every instance we can run, including open-shell multireference oxides (CrO, NiO) and EUV Sn-oxides; (ii) scalability — QSCI reaches chemical accuracy on a vanishing fraction of the CI space (3.8% at 20q, 0.15% at 28q), with pool compression holding accuracy at 25–40% of the doubles pool; (iii) transferable structure — a generator trained on small systems proposes good circuits for a larger one it never saw (§3). Honest boundary: at sizes classical FCI/DMRG still solve these instances we show correctness and favourable scaling, not a head-to-head speed win — that win is the 40q+ regime the GPU run targets (Table 2).

Where it matters most: trust under strong correlation. Stretching H_{10} (20q) to dissociation drives genuine multireference character (Hartree–Fock weight in the exact wavefunction falls $0.93 \rightarrow 0.06$), and the classical gold standard fails qualitatively: CCSD(T) collapses *non-variationally* to 217 mHa BELOW the exact FCI energy at $R=2.4 \text{ \AA}$ — a confidently wrong, unphysical answer. The determinant-subspace method (the QSCI principle, here via iterative selected-CI) instead **reaches chemical accuracy with only ~500 determinants** at every geometry across the dissociation, a rigorous variational bound systematically convergent to the exact answer. **Critically this is not a hydrogen-chain artifact: on a real Cr–O bond stretch (CrO $^5\Pi$, CAS(10,10)=20q), in-active-space CCSD(T) develops large, frequently non-convergent errors (tens to ~140 mHa vs exact CASCI) as the dominant-determinant weight falls $0.87 \rightarrow 0.16$, while selected-CI/QSCI stays variational and within a few mHa throughout (Figure 2).** And we certify the error rather than assert it: an Epstein–Nesbet PT2 correction tightens the rigorous variational upper bound toward FCI — **E_var+PT2 is a second-order estimate (not a bound) that at our budgets converges to FCI from above:** near equilibrium the CIPSI extrapolation reproduces FCI to ~4 mHa ($R^2=0.999$), and even at the strongest correlation ($R=2.4 \text{ \AA}$, 468 determinants) E_var+PT2 sits +10.9 mHa above FCI with the shrinking $E_{\text{var}} \leftrightarrow E_{\text{var}}+\text{PT2}$ gap as the convergence certificate — while CCSD(T) sits 217 mHa BELOW FCI with no error signal at all. (Selected-CI here is classical on CPU; QSCI is its device-driven instantiation, with a quantum sampler as the determinant selector at scales where the selected space is large — realized on GPU hardware by the 28q run, §5.) **R&D value:** a confidently-wrong ranking (a flipped spin state, a non-variational binding energy) commits multi-month synthesis spend to a false lead — the dominant hidden cost of discovery, which a self-certifying variational selector that stays correct where DFT/CCSD(T) silently fail directly reduces.




Figure 2. Trust on a real oxide. As the CrO $^5\Pi$ bond stretches (stronger multireference character, left $\rightarrow$ right), in-active-space CCSD(T) becomes erratic and frequently non-convergent (open markers), a few to ~140 mHa vs exact CASCI, while selected-CI/QSCI stays a variational bound within the chemical-accuracy band (≤ 2.8 mHa even at strongest correlation). CCSD(T) uses the identical embedded active-space Hamiltonian as CASCI (apples-to-apples).

Frozen extension campaign (E2–E5; pre-registered before access, executed 2026-07-23, reported as measured). Five protocols were git-timestamped in preregistration_v2.json before access; reported exactly as measured — pass, fail, or resource-DNF. **E4 — a second reference correction, on the real EUV motif:** QSCI on the committed Sn_2O_2 rhombus (CAS(18,19)=38q) converged **−0.399 mHa BELOW the same-CAS DMRG($\chi=400$) reference** (524,764 determinants, 7.4 h; reference committed pre-run) — the CrO reference-correction reproduced on the actual tin-oxo chemistry, not system-specific. **EUV-motif trust curve:** stretching the Sn_2O_2 Sn–O bridge 2.05 → 3.28 Å (the strongly-correlated cleavage regime EUV resist chemistry occupies), in-active-space CCSD(T) error grows **0.14 → 5.49 mHa (~40×)** while the identical selected-CI/QSCI stays variational and accurate (≤ 0.47 mHa everywhere) and the dominant-determinant weight collapses 0.95 → 0.53 — Figure 2's trust story, now on the real motif under bond cleavage. **E3 — the 40q absolute-certification protocol (terminal at it5):** growth-to-PT2-certificate on H_{20} 40q with a bucket-parallel Epstein–Nesbet estimator (bit-identical to serial); the run was ended by an external pod kill during it6 growth (disclosed operational outcome, not a convergence failure). At its it5 terminal: **E_var +0.185 mHa vs $\chi=400$ (prediction ii MET), 750,257 dets (iii MET),** the |PT2| certificate a clean converging trace 84 → 4.6 → 2.6 → 2.0 → 1.57 → 1.31 mHa; the |PT2| ≤ 0.5 absolute point (i, $\sim$ it10–11) unreached, reported as-measured. **E2 — resource-DNF, disclosed as an operational record (not the E2 result):** the device-seeded 40q equivalence run (committed 102-determinant seed) needs ≥ 128 GiB and was OOM-killed by the cgroup at iter-3 cap-fill (peak 133.5 GB, 97% of limit; two deterministic attempts). E4 completed on the same box at 59.7 GB — the +2 qubits double the footprint. An honest hardware ceiling, no tuning. **E5 — the least-grounded frontier (44q), non-converged and reported as such:** H_{22} 44q growth against a re-frozen $\chi=1200$ judge, terminated early to protect E3's runway; committed trajectory +125 → +8.7 → +5.2 → +4.0 mHa (decelerating), not reaching the ≤ 1.6 mHa threshold before the abort gate — the pre-registered expected outcome, no threshold moved and frozen parameters untouched.

Proxy validated against real measurement; cost scaling. We confirm the determinant-space proxy tracks the real circuit-sampled pipeline by executing measured QSCI (qml.sample, 3,000 shots $\times$ 160 circuits) at 16q and 20q — measured energies match the proxy on the trained generator (16q: 28.1 vs 26.6; 20q: 50.4 vs 49.3 mHa), extending the sampled pipeline past 12q (at this small shot budget diffuse random sampling edges out the generator on absolute energy — a low-shot coverage effect, not a quality reversal). Determinants for chemical accuracy grow $\sim 1.38\times$ per qubit vs $2.0\times$ for full FCI (log $R^2=0.99$) — a vanishing fraction but still exponential ($\sim 10^6$ dets at 40q+); we do not claim polynomial scaling.

6. Platform Use and Resourcing

qBraid provides classical (CPU/GPU) and quantum (QPU) credits. We use **NVIDIA H100/A100 (80 GB)** with CUDA-Q: tensornet-mps for the 24–40q tier, cuStateVec for exact validation to ~ 32 q; 4–8 GPUs (NVLink) enable distributed evaluation (pillar 4) and the >40q bonus; QPU (IonQ/IBM) for 10–16q validation. The ~ 16 TB exact statevector at 40q collapses onto one 80 GB GPU via CUDA-Q MPS ($\sim 200\times$ reduction) — what makes the 40-qubit tier reachable.

Resource estimate. UCC excitations decompose to ≤ 2 -qubit gates; with pool compression a length-8 circuit is shallow, within MPS reach near equilibrium. QSCI needs $\sim 10^3$ – 10^4 shots/circuit. We estimate ~ 1 – 2 weeks single-GPU wall-clock for the full 24 → 40q sweep plus oxide runs. Measured: 20q 191 s, 28q 2,386 s total — GPU sampling 1–32 s, the rest classical determinant growth (the wall at scale, cut $\approx 5\times$ by the incremental engine).

7. Limitations and Honest Scope

We state where the approach does and does not yet provide benefit. (i) **Measurement vs proxy:** the GQE → QSCI loop is measured at 12q and executed with real GPU sampling at 20q/28q (§5); larger QSCI numbers use a hardware-independent determinant-space proxy. (ii) **No classical-beating claim yet:** at ≤ 28 q classical FCI/DMRG already solve these instances, so we show correctness and favourable scaling, not a head-to-head speed advantage; that regime is 40q+ strong correlation where DMRG bond dimension explodes. (iii) **40q absolute certification:** the flagship meets its pre-registered relative criterion; the frozen E3 certificate reached a terminal it5 (external pod kill during it6 growth), prediction ii met, the |PT2| ≤ 0.5 absolute point unreached — reported as-measured (§5e). (iv) **Transfer scope:** cross-size transfer is a 3-seed-mean trend narrowing into noise beyond ~ 28 q; cross-chemistry wins 3 of 6 oxides (BeO a tie), fails on SnO; target-specific MP2 beats the transferred prior — a zero-cost prior, not a universal or MP2-beating one. (v) **Strong correlation:** where CCSD(T) collapses (§5: H_{10} , real CrO and Sn_2O_2), the determinant-subspace method stays a rigorous variational bound (EN-PT2 tightening from above), but strict accuracy at the strongest correlation needs the large-subspace device-sampled regime.

8. Conclusion and Reproducibility

MATGEN-Q is a working two-stage GQE — tensor-network + QSCI tiers with MP2 pool compression — targeting 40 qubits on a single GPU, every claim third-party reproducible via a one-command *reproduce.py* and a “Launch on qBraid” README that let judges re-run each headline result unmodified.

References

- [1] K. Nakaji, et al. “The generative quantum eigensolver (GQE) and its application for ground state search.” arXiv:2401.09253 (2024). [AIST]
- [2] N. P. D. Sawaya, et al. “HamLib: A library of Hamiltonians for benchmarking quantum algorithms and hardware.” Quantum 8, 1559 (2024).
- [3] S. Minami, K. Nakaji, et al. “Generative quantum combinatorial optimization ... conditional generative quantum eigensolver.” Digital Discovery 4(8) (2025).
- [4] K. Kanno, et al. “Quantum-selected configuration interaction ... subspaces selected by quantum computers.” arXiv:2302.11320 (2023).
- [5] NVIDIA Corporation. CUDA-Q and cuQuantum SDK (cuStateVec, tensornet / tensornet-mps backends).
- [6] J. Tilly, et al. “The Variational Quantum Eigensolver: A review of methods and best practices.” Physics Reports 986 (2022).
- [7] C. Fare, et al. “A multi-fidelity machine learning approach to high-throughput materials screening.” npj Computational Materials 8, 257 (2022).
- [8] T. D. Kharazi, et al. “Quantum Simulations for Extreme Ultraviolet Photolithography.” arXiv:2602.20234 (2026). [Xanadu & Mitsubishi Chemical]
- [9] R. Kemmoku, Q. Gao, S. Kanno, et al. “Generative Circuit Design for Quantum-Selected Configuration Interaction.” arXiv:2604.09756 (2026). [Mitsubishi Chemical]