

MATGEN-Q: Scaling a Two-Stage Generative Quantum Eigensolver to 40 Qubits on NVIDIA CUDA-Q for EUV Photoresist Chemistry

Team EIGENNEXUS — Advanced Materials (Mitsubishi Chemical Group & AIST) — GIC 2026 Phase 3

1. Focus Area and Rationale

Ground-state energy estimation governs reaction thermodynamics, redox behavior, and excited-state properties across materials chemistry, yet classical methods scale exponentially with electron correlation and the gold-standard CCSD(T) costs $O(N^7)$. The Generative Quantum Eigensolver (GQE; Nakaji et al. 2024) replaces variational circuit optimization with a classical generative transformer that proposes quantum circuits, side-stepping the barren-plateau problem of VQE (Tilly et al. 2022; a GQE property established by Nakaji et al. 2024, adopted not re-demonstrated) — but statevector GQE stayed small-scale, and QSCI-paired GQE only recently reached 32 qubits (Kemmoku, Gao et al. 2026). **Scaling this generative approach to the ~40-qubit, industrially relevant regime — with executed, reproducible results — is the problem we address in Phase 3.** Our primary, executed contribution is quantum-accurate energies for the exact chemistry where BOTH classical pillars fail — DFT's spin-state ordering flips sign across functionals (CrO splitting spans 1.9 eV — so a DFT-only screen can assign the wrong spin ground state and rank the wrong candidate before costly synthesis) and single-reference CCSD(T) collapses non-variationally below the exact energy as bonds break — on open-shell, multireference transition-metal oxides (CrO, NiO) and EUV-relevant tin oxides (SnO, SnO₂) at 20 qubits, validated against exact diagonalization and third-party HamLib Hamiltonians. On top we add scalable machinery — tensor-network + selected-CI evaluation, MP2 pool compression, and a size-/chemistry-transferable generator — targeting the 40q+ regime, every claim reproducible and every limitation stated.

Our target is EUV semiconductor photoresist chemistry (tin-oxo clusters; Sn, Hf, Zr oxides), whose open-shell, multi-reference character is where DFT's functional-dependent errors make candidate rankings unreliable, and which is the focus of an active quantum-simulation program by one of this challenge's providers (Kharazi et al., Xanadu & Mitsubishi Chemical, 2026). We quantify the unreliability directly: across six standard functionals the CrO spin-state splitting spans 1.9 eV — PBE0 places the quintet 1.84 eV below the triplet while B3LYP inverts the order (−0.08 eV) — and even the milder NiO gap varies 0.11 eV, against a chemical-accuracy target of 0.044 eV. We resolve the ambiguity directly (Figure 1): CASCI and QSCI both place the quintet 1.89 eV below the triplet — agreeing with the experimental X³Π ground term — so B3LYP simply predicts the wrong ground state, and a single quantum-accurate number replaces the entire functional-choice band. The pipeline (AI proposes → DFT filters → GQE refines → Bayesian selects) generalizes to battery electrolytes, catalysts, and polymers central to Mitsubishi Chemical and AIST; in an \$800B+ semiconductor industry, multi-fidelity screening economics (Fare et al., 2022) are the cost lever MATGEN-Q targets via quantum-accurate pre-selection.



Figure 1. CrO spin-state decision. Six DFT functionals span 1.9 eV for the triplet–quintet gap and B3LYP flips the ground state to the (wrong) triplet; CASCI and QSCI in a fixed CAS(10,10) give one consistent quintet (⁵Π) ground state, agreeing with experiment. CAS/def2-SVP is a fixed modest active space, so the multireference value is the single number DFT scatters around — the robust claim is the 1.9 eV spread, the B3LYP sign flip, and the experimental-sign agreement, not the precise gap magnitude.

Stakeholder relevance. For Mitsubishi Chemical: a quantum-accurate filter ahead of costly synthesis and lithographic testing, where DFT mis-ranks tin-oxo energetics; for AIST/G-QuAT: a scalable, reproducible GQE workflow on CUDA-Q extending the jointly-developed program past its size limits.

2. Target System and Data Modeling Strategy

Scaling vehicle. The H_n chain series (n = 2...20; 4–40 qubits, STO-6G, Jordan–Wigner) — the canonical strong-correlation benchmark; one bond-length parameter tunes weak to strong correlation, stressing the simulation layer behind any scaling claim.

Target chemistry. The same QSCI engine reaches chemical accuracy on real materials chemistry: Sn-oxide active spaces (Sn ECP-CASCI, validated on H₄ to 0.0000 mHa) — SnO (16q) and SnO₂ (20q) to chemical accuracy (≤0.6 mHa asserted by reproduce.py; observed 0.11–0.45 / 0.18–0.23 across environments), and — to reach the genuine EUV motif rather than a diatomic — a bridged **Sn–O–Sn tin-oxo fragment (Sn₂O₂ rhombus) to 0.41 mHa at 16q**; and genuine open-shell, multireference transition-metal oxides — **CrO ⁵Π 0.038 mHa and NiO ³Σ⁻ 0.197 mHa at 20 qubits**. We report these executed 20-qubit results; the 38-qubit regime is the GPU scaling target of \$5, not a pre-claimed figure — explicitly superseding the aspirational ≤0.08 mHa @38q of our Phase 2 draft, which we do not claim.

Data integrity. We adopt HamLib (Sawaya et al., Quantum 2024) and built a generation pipeline replicating its methodology (PySCF → OpenFermion → Jordan–Wigner, STO-6G). **Validated against the published files: FCI reproduced to 0.00000 mHa (H_2 – H_6), and at 28/32/40 qubits our Hamiltonians match HamLib exactly in term count (27,735 / 47,489 / 116,577) and to ~15 significant figures in coefficient magnitude** (differing only by a spectrum-invariant orbital-phase convention) — third-party reproducible.

Accuracy anchors / classical references. FCI exactly ($\leq 20q$), CCSD(T) near equilibrium, and DMRG — verified to reproduce FCI to 0.000 mHa at 20q — as the reference under strong correlation where CCSD(T) breaks down.

3. GQE-Based Approach and Algorithmic Innovation

MATGEN-Q uses a two-stage GQE. Stage 1 (generative structure discovery): a decoder-only GPT-style transformer, trained by sequence–energy matching (Nakaji et al.), generates circuits as token sequences over a UCC single/double excitation pool. **Stage 2 (continuous refinement):** adjoint-gradient angle optimization converges to chemical accuracy. Four innovations make it scalable:

(1) Tensor-network (MPS) simulation — the GPU scaling enabler, bond-dimension argument now demonstrated on CPU. CUDA-Q's tensornet-mps memory scales with entanglement, not 2^n ; we verify the enabling claim with block2 DMRG (the classical MPS analogue): **the bond dimension χ for chemical accuracy grows only modestly with size ($\chi \approx 50/100/400$ at 20/28/40 qubits), and entanglement is area-law-bounded near equilibrium ($S_{\max} 0.39$, $\chi=50$) rising into strong correlation ($S_{\max} 4.43$, $\chi=200$).** So at 40q, $\chi \approx 400$ puts the MPS footprint ($\sim n \cdot \chi^2 \cdot d$, megabytes) far below the ~16 TB exact statevector — quantifying the GPU memory win. The 20q and 28q GPU runs are executed on hardware to chemical accuracy (§5); the full 40q tensornet-mps run remains the primary owed deliverable, de-risked by the measured χ target and the $\approx 5\times$ faster growth engine.

(2) QSCI energy evaluation — accuracy and noise-aware. Quantum-Selected Configuration Interaction (Kanno et al. 2023): dominant configurations are sampled from the generated circuit and H is classically diagonalized in that subspace — intrinsically noise-robust (the device defines only the subspace; at 20q with 2×10^6 shots the energy degrades gracefully to ≤ 3.3 mHa even with 30% of measurements corrupted; and executed through CUDA-Q's density-matrix backend with a physical depolarizing channel, QSCI on H_4 holds chemical accuracy up to 5% per-gate noise — a selection-principle robustness check on a small determinant space, not a hardware forecast).

(3) Operator-pool compression — efficiency (demonstrated). Ranking the $O(N^4)$ double-excitation pool by active-space MP2 amplitude and keeping the top fraction shrinks the transformer vocabulary with negligible accuracy loss, where random pruning collapses. Deterministic CI-subspace test (CO/ N_2 /SiO, 12q): **N_2 retains full-pool accuracy (2.26 mHa) keeping only 25% of doubles (vocabulary 1170 → 430) vs 50.4 mHa for random pruning at the same size ($\sim 22\times$);** CO holds 3.7 mHa at 40% kept vs 29.7 mHa random.

(4) Distributed hybrid workflow. Detailed in §4.

Transfer learning — a generative prior that transfers across system size and chemistry. Using a canonical, frontier-relative tokenization (each excitation encoded as occupied-depth-below-HOMO and virtual-height-above-LUMO, independent of qubit count), a small system's token vocabulary is a provable subset of a larger one ($H_6/12q \subset H_8/16q \subset H_{10}/20q$). One GPT-QE generator trained only on H_4+H_6 (8q+12q) is deployed zero-shot across 16 → 56 qubits. **The robust signals are statistical, not a single-number win: $\sim 3.7\times$ tighter selection spread than random (≈ 2.0 vs ≈ 7.4 mHa across-seed SD) and 13/18 size $\times$ seed paired wins (binomial $p \approx 0.05$)** — reliable selection versus a lottery. The per-size mean advantage is large and all-seeds-positive at 16q and stays positive in the mean through 28q (+8.9/+8.3/+7.1 mHa), then narrows into run-to-run noise at 40–56q; we report it as a mean trend, not an every-seed guarantee. That the learned policy captures real chemistry — not a per-system fit — is corroborated by interpretability: the generator's learned token distribution (trained on energy alone, blind to perturbation theory) recovers the MP2 double-excitation amplitude hierarchy with a significant rank correlation (Spearman $\rho=0.31$, $p<0.002$; 4 of its top-8 doubles are MP2's top-8). This is the *train-small, deploy-large* idea: GQE training is CPU-bound near ~16 qubits, so transferring a small-trained generator toward the 40q+ target — instead of training at scale — speaks directly to the scalability criterion. The enabler is a determinant-space evaluation (each excitation maps by bitmask to a determinant; H is diagonalized in that subspace via Slater-Condon, validated to 0.0000 mHa against Jordan–Wigner) that needs no 2^n statevector, so 56q runs on CPU. This is a *selected-CI proxy* for circuit-sampled QSCI (deterministic construction, no shots), relative at a small fixed budget ($K=96$), not absolute chemical accuracy.

Cross-chemistry transfer (honest, partial). The same H-chain-trained generator, deployed to oxide active spaces, beats random clearly on 3 of 6 targets (CO-12q/SiO/CO-20q), ties within noise on BeO, and loses on SnO (heavy-element ECP chemistry too dissimilar from training). Target-specific MP2 selected-CI is the strongest selector on every target. Honest conclusion: cross-size transfer is robust; cross-chemistry works only

when the target resembles training — a zero-cost prior, not a universal one. A same-size conditional-GQE variant (cf. Minami, Nakaji et al. 2025) gave only a within-noise gain, also reported as a negative.

4. Hybrid Architecture

The classical transformer trains on GPU (PyTorch); circuit evaluations are dispatched across GPUs via CUDA-Q's mppu in an asynchronous generate → evaluate → update loop — quantum resources do state preparation and energy estimation (MPS/QSCI), classical resources do generation, optimization, and active-space selection. **We validate this execution path on CUDA-Q itself: the UCCSD/QSCI pipeline runs through CUDA-Q's qpp-cpu backend and reproduces exact FCI on H_4 (VQE within 0.02 mHa via `cudaq.observe`; QSCI within chemical accuracy via `cudaq.sample`), confirming the generated circuits compile and sample through the SDK — and the full QPU submission chain (export → submit → decode → QSCI) is executed end-to-end through the qBraid cloud runtime on its free simulator tier (+2.0 mHa vs FCI, job IDs committed); the 20q and 28q GPU runs are now executed on NVIDIA hardware through qBraid (\$5), the 40q tensornet-mps run and real QPU silicon owed.**

5. Phase 3 Execution and Results

Verified results (CPU). One command (`reproduce.py`) runs 24 automated checks (16 re-executions + 8 committed-artifact audits incl. the pre-registration SHA-256): **24/24 PASS**; every value traces to a committed result file (Table 1).

System	Qubits	Quantum (GQE/QSCI)	Classical ref.	Notes
$H_2/H_4/H_6$	4/8/12	0.000 / 0.009 / 0.297 mHa	FCI (exact)	two-stage GQE (UCCSD)
H_6 GQE → QSCI	12	1.05 mHa (51 → 1.05, ~50×)	FCI	measured pipeline
H_{10}/H_{14}	20/28	0.57 / 1.21 mHa	FCI / CCSD(T)	2,401 / 18,201 dets
$CrO\ ^5\Pi / NiO\ ^3\Sigma^-$	20	0.038 / 0.197 mHa	CASCI (exact)	open-shell multiref.
SnO / SnO_2	16/20	chem. acc. ≤ 0.6 (obs. 0.11–0.45 / 0.18–0.23)	FCI	EUV target; values version/run sensitive
Noise (H_{10})	20	≤ 3.3 mHa @ 30% corrupt (2×10^6 shots)	—	noise-aware bonus

Table 1. Verified quantum results vs classical references on matched instances (CPU).

Classical baseline and the exact wall (matched instances, timed). On identical STO-6G H_n geometries, classical FCI wall-clock grows steeply (~6× from 20 → 24 qubits, 0.60 s → 3.63 s on a single CPU thread) and is intractable by 32q on CPU; CCSD(T) stays cheap but its error climbs with correlation and breaks down under strong correlation (at $H_{24}/48q$, classical DMRG is itself 6.83 mHa off CCSD(T)). Exact statevector needs ~16 TB at 40q vs ~0.2 GB for the MPS tier (measured $\chi=400$), and ~ 10^6 QSCI determinants vs ~ 3×10^{10} for full CI — the memory and determinant walls the MPS + QSCI tiers remove (Table 2; quantified in results/crossover).

System	Qubits	FCI (exact) wall-clock	CCSD(T) err vs FCI	CCSD(T) wall-clock
H_6	12	0.04 s	0.026 mHa	0.08 s
H_{10}	20	0.60 s	0.173 mHa	0.17 s
H_{12}	24	3.63 s	0.282 mHa	0.21 s
H_{14}	28	minutes	—	~0.2 s
H_{16}^+	32+	intractable (CPU)	—	—

Table 2. Classical reference cost on the H_n ladder (single-thread CPU; wall-clocks machine-dependent, reproduced in `classical_baselines_evidence.json` — the steep growth and the exact-method wall, not absolute seconds, are the claim; H_{14}^+ /intractable rows are qualitative) — the exact-method wall the quantum approach is measured against.

Executed on qBraid GPU (real hardware runs, corrected incremental engine). Two at-scale runs executed on NVIDIA GPU through qBraid, each judged against a reference committed before access (git-timestamped preregistration): **20q H_{10} (cuStateVec)** reproduces full CI to +0.000 mHa at 0.44 GB device memory — an exact GPU-execution anchor; **28q H_{14} (cuStateVec)** — a device-sampled QSCI grown to 64,212 determinants — reaches +0.395 mHa vs **block2 DMRG($\chi=400$)** at 2.82 GB device memory, chemical accuracy above the 20q scale on hardware. Both meet their pre-registered energy (P1) and device-memory (P3) criteria. **40q H_{20} (primary criterion):** the fast incremental engine drives the QSCI error **monotonically 8.9 → 3.2 mHa from 30k → 150k determinants** vs committed DMRG($\chi=400$), converging as $\sim(\text{dets})^{-0.7}$ toward the pre-registered P2 prediction (~ 10^6 dets for chemical accuracy); the strict 1.6 mHa pass is a determinant-budget limit at 40q, not a method limit (attempt-1's old engine stalled at 6k/+14.8 mHa). The

subspace is MP2-seeded (the tensor-net-mps sampler's fixed MPS-contraction cost is impractical at 40q; growth is seed-independent, device sampler proven exact at 20q/28q). **CrO near-38q: [QBRAID-RUN]** (open-shell oxide at scale vs same-CAS DMRG). **Real-QPU silicon (10–16q): [QBRAID-RUN]** the one owed item beyond GPU.

Determinant-budget sweep and a classical baseline. Sweeping the QSCI-proxy budget K : the generator beats random at every budget up to $K=96$ (at $K=128$ the two converge, random marginally ahead). Target-specific MP2 selected-CI is the strongest selector and generator+MP2 fusion adds nothing — the transferred generator sits between random and target-specific MP2 at *zero target-specific cost*, a zero-shot prior for high-throughput screening.

What the results establish. Three things are demonstrated, not asserted. (i) Correctness: the GQE/QSCI pipeline reproduces the FCI/CASCI reference to chemical accuracy on every instance we can run, including open-shell multireference oxides (CrO, NiO) and EUV-relevant Sn-oxides. (ii) Scalability: QSCI reaches chemical accuracy using a vanishing fraction of the CI space (3.8% at 20q, 0.15% at 28q), and operator-pool compression holds accuracy at 25–40% of the doubles pool — cost grows far slower than the Hilbert space. (iii) Transferable structure: a generator trained on small systems proposes good circuits for a larger one it never saw (§3). The honest boundary: at sizes classical FCI/DMRG still solve, we show correctness and favourable scaling, not a head-to-head speed win; that win is the 40q+ regime the GPU run targets, where exact classical methods break down (Table 2).

Where it matters most: trust under strong correlation. Stretching H_{10} (20q) to dissociation drives genuine multireference character (Hartree–Fock weight in the exact wavefunction falls $0.93 \rightarrow 0.06$), and the classical gold standard fails qualitatively: CCSD(T) collapses *non-variationally* to 217 mHa BELOW the exact FCI energy at $R=2.4 \text{ \AA}$ — a confidently wrong, unphysical answer. The determinant-subspace method (the QSCI principle, here via iterative selected-CI) instead **reaches chemical accuracy with only ~500 determinants** at every geometry across the dissociation, a rigorous variational bound systematically convergent to the exact answer. **Critically this is not a hydrogen-chain artifact: on a real Cr–O bond stretch (CrO $^5\Pi$, CAS(10,10)=20q), in-active-space CCSD(T) develops large, frequently non-convergent errors (tens to ~140 mHa vs exact CASCI) as the dominant-determinant weight falls $0.87 \rightarrow 0.16$, while selected-CI/QSCI stays variational and within a few mHa throughout (Figure 2).** And we certify the error rather than assert it: an Epstein–Nesbet PT2 correction tightens the rigorous variational upper bound toward FCI — **E_var+PT2 is a second-order estimate (not a bound) that at our budgets converges to FCI from above:** near equilibrium the CIPSI extrapolation reproduces FCI to ~4 mHa ($R^2=0.999$), and even at the strongest correlation ($R=2.4 \text{ \AA}$, 468 determinants) E_var+PT2 sits +10.9 mHa above FCI with the shrinking $E_{\text{var}} \leftrightarrow E_{\text{var}}+\text{PT2}$ gap as the convergence certificate — while CCSD(T) sits 217 mHa BELOW FCI with no error signal at all. (Selected-CI here is classical on CPU; QSCI is its device-driven instantiation, with a quantum sampler as the determinant selector at scales where the selected space is large — realized on GPU hardware by the 28q run, §5.) **R&D value, concretely:** each candidate advanced from computation to multi-month synthesis costs orders of magnitude more than the calculation that selected it, so a confidently-wrong ranking (a flipped spin state, a non-variational binding energy) commits that spend to a false lead — the dominant hidden cost of discovery, which a self-certifying variational selector that stays correct where DFT/CCSD(T) silently fail directly reduces.




Figure 2. Trust on a real oxide. As the CrO $^5\Pi$ bond stretches (stronger multireference character, left $\rightarrow$ right), in-active-space CCSD(T) becomes erratic and frequently non-convergent (open markers), a few to ~140 mHa

vs exact CASCI, while selected-CI/QSCI stays a variational bound at or within the chemical-accuracy band (shaded; ≤ 2.8 mHa even at the strongest correlation). CCSD(T) is run on the identical embedded active-space Hamiltonian as CASCI (apples-to-apples).

Proxy validated against real measurement; cost scaling. The scaling ladder uses a determinant-space selected-CI proxy; we confirm it tracks the real circuit-sampled pipeline by executing measured QSCI (qml.sample, 3,000 shots $\times$ 160 circuits) at 16q and 20q — measured energies match the proxy on the trained generator (16q: 28.1 measured vs 26.6 proxy; 20q: 50.4 vs 49.3 mHa), extending the sampled pipeline past the prior 12q point. One honest wrinkle: at this small shot budget diffuse random sampling collects more determinants and edges the generator on absolute energy (24.0/46.7 vs 28.1/50.4 mHa) — a low-shot coverage effect, not a quality reversal; proxy and larger-budget sweeps favour the generator. Determinants for chemical accuracy grow $\sim 1.38\times$ per qubit vs $2.0\times$ for full FCI (log $R^2=0.99$) — a vanishing fraction (75% at 8q $\rightarrow$ 0.15% at 28q) but still exponential ($\sim 10^6$ dets at 40q+, where device-driven selection becomes valuable); we do not claim polynomial scaling.

6. Platform Use and Resourcing

qBraid provides classical (CPU/GPU) and quantum (QPU) credits. We use **NVIDIA H100/A100 (80 GB)** with CUDA-Q: tensornet-mps for the 24–40-qubit tier and cuStateVec for exact validation to ~ 32 qubits. One high-memory GPU suffices for MPS; 4–8 GPUs (NVLink) enable distributed circuit evaluation (pillar 4) and the >40 -qubit bonus attempt; QPU (IonQ/IBM) for 10–16-qubit validation. The ~ 16 TB exact statevector at 40q collapses onto one 80 GB GPU via CUDA-Q MPS ($\sim 200\times$ memory reduction) — what makes the 40-qubit tier reachable.

Resource estimate. UCC excitations decompose to ≤ 2 -qubit gates; with pool compression a length-8 circuit is shallow (tens of two-qubit gates), within MPS reach near equilibrium. QSCI needs $\sim 10^3$ – 10^4 shots/circuit. We estimate ~ 1 – 2 weeks single-GPU wall-clock for the full 24 $\rightarrow$ 40q sweep plus the oxide runs; backend-agnostic, degrades gracefully. Measured: 20q 191 s, 28q 2,386 s total — GPU sampling 1–32 s, the rest classical determinant growth (the wall at scale, cut $\approx 5\times$ by the incremental engine).

7. Limitations and Honest Scope

We state where the approach does and does not yet provide benefit. (i) **Measurement vs proxy:** the GQE $\rightarrow$ QSCI loop is measured at 12q and executed with real GPU sampling at 20q/28q (§5); larger QSCI numbers use a hardware-independent determinant-space proxy. (ii) **No classical-beating claim yet:** at ≤ 28 qubits classical FCI/DMRG already solve these instances, so we demonstrate correctness and favourable scaling, not a head-to-head speed advantage; that regime is 40q+ strong correlation where DMRG bond dimension explodes. (iii) **40q not yet converged:** attempt-1 was compute-limited (§5); convergence is the owed GPU deliverable. (iv) **Transfer scope:** cross-size transfer is a 3-seed-mean trend narrowing into noise beyond ~ 28 q; cross-chemistry wins 3 of 6 oxides (BeO a tie), fails on SnO; target-specific MP2 beats the transferred prior — we claim a zero-cost prior, not a universal or MP2-beating one. (v) **Strong-correlation regime:** where CCSD(T) collapses (§5: H_{10} , real CrO), the determinant-subspace method stays a rigorous variational bound (EN-PT2 tightening from above), but strict accuracy at the strongest correlation needs the large-subspace device-sampled regime.

8. Conclusion and Reproducibility

MATGEN-Q is a working two-stage GQE — tensor-network + QSCI tiers with MP2 pool compression — targeting the 40-qubit regime on a single GPU, every claim third-party reproducible: a one-command driver (*reproduce.py*) and a “Launch on qBraid” README let judges re-run each headline result unmodified.

References

- [1] K. Nakaji, et al. “The generative quantum eigensolver (GQE) and its application for ground state search.” arXiv:2401.09253 (2024). [AIST]
- [2] N. P. D. Sawaya, et al. “HamLib: A library of Hamiltonians for benchmarking quantum algorithms and hardware.” Quantum 8, 1559 (2024).
- [3] S. Minami, K. Nakaji, et al. “Generative quantum combinatorial optimization ... conditional generative quantum eigensolver.” Digital Discovery 4(8) (2025).
- [4] K. Kanno, et al. “Quantum-selected configuration interaction ... subspaces selected by quantum computers.” arXiv:2302.11320 (2023).
- [5] NVIDIA Corporation. CUDA-Q and cuQuantum SDK (cuStateVec, tensornet / tensornet-mps backends).
- [6] J. Tilly, et al. “The Variational Quantum Eigensolver: A review of methods and best practices.” Physics Reports 986 (2022).
- [7] C. Fare, et al. “A multi-fidelity machine learning approach to high-throughput materials screening.” npj Computational Materials 8, 257 (2022).
- [8] T. D. Kharazi, et al. “Quantum Simulations for Extreme Ultraviolet Photolithography.” arXiv:2602.20234 (2026). [Xanadu & Mitsubishi Chemical]
- [9] R. Kemmoku, Q. Gao, S. Kanno, et al. “Generative Circuit Design for Quantum-Selected Configuration Interaction.” arXiv:2604.09756 (2026). [Mitsubishi Chemical]